

Death Rate Correlation with COVID-19 X-Ray Images

Overview

Team Members

1. Harshal Sharma

I'm currently pursuing a BS in CS at the Texas Tech University.
(<https://www.linkedin.com/in/harshalsharma08>)

2. Shreyash Shrivastava

I am an incoming Analyst at Bank of America in Plano, TX. I have completed a BSCS from UT Arlington.
(<https://www.linkedin.com/in/shreyash-shrivastava/>)

Project Development libraries

Numpy, Pandas, Matplotlib, Jupyter notebook, Seaborn, PIL, Keras, OpenCV, Scikit-learn

Dataset Source

1. <https://www.kaggle.com/bachrr/covid-chest-xray/data>
2. <https://www.kaggle.com/tawsifurrahman/covid19-radiography-database>

Problem Statement

To find a correlation between chest x-ray images of COVID-19 infected patients and their subsequent mortality.

Procedure and Efficacy

Dataset Explanation:

Dataset1

We have labeled dataset of COVID-19 infected patient's chest X-ray depicting the mortality label for each image. In the same set, we also have unlabeled images of chest X-rays as well.

Dataset2

We have three sets of folders, namely; Viral Pneumonia, COVID-19, and Normal. The three folders consist of chest X-rays of people with the said conditions.

CNN (Convolution Neural Network):

1. Clean and preprocess the raw data using Keras preprocessing library. Separate the data with missing values. Then, create a CNN model as per the described architecture.
2. Use the clean data to train the model with specified hyperparameters.
3. The model calculates the chance of survival from the input images and classifies it into two classes: *WILL_SURVIVE*, *WILL_NOT_SURVIVE*
4. Use out of setting data (Dataset 2) which consists of three categories: *COVID-19*, *NORMAL*, *VIRAL PNEUMONIA*
5. Clean and preprocess all three categorical data, and use the trained model to classify it. Then, based on the classification calculate the death rate and survival rate.
6. (Next steps) Use GAN to generate artificial images and use that to again train the model and evaluate it.

Constraints

Limited Data available

Limited Computational power

Time constraint

Results

The CNN model performed extremely well on online datasets for COVID-19 patients. The results were astonishingly close to the real world cases despite the computational and data scarcity constraints we possessed.

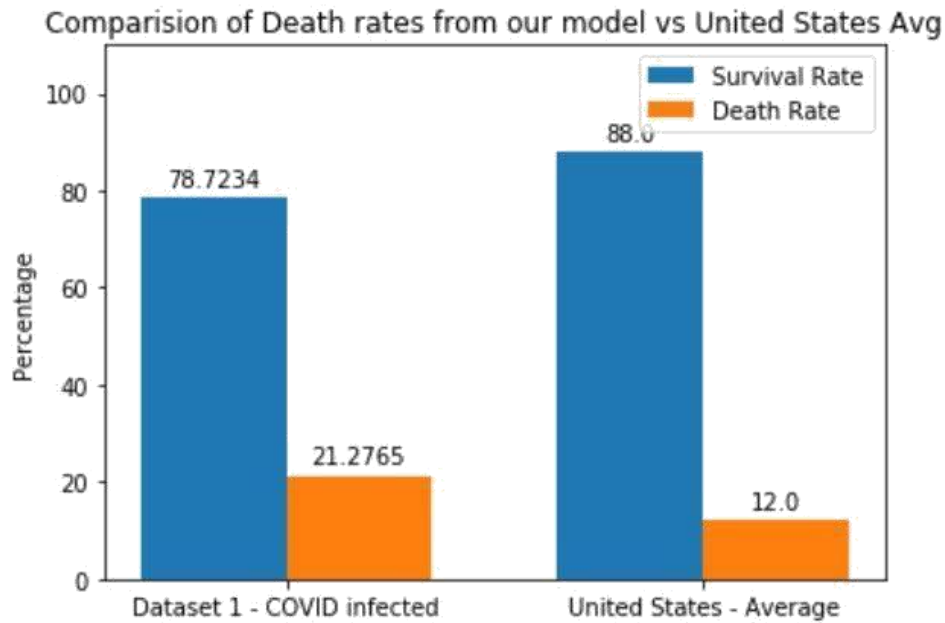


Figure 1 - Comparison of Death rate prediction of our model vs the United States closed cases

NOTE:

The definition of the real-world death rate should be pondered on. The death rate should be calculated on the number of closed cases. This implies that the on-going cases of COVID-19 infection should be deliberately excluded from the calculation. The closed cases signify two possibilities, either the infection recovers or its consequences into a fatality.

Therefore, the death rate is calculated as $(\text{Total Deaths}) / (\text{Total number of people recovered} + \text{Total Deaths})$

Comparison of Death rates from our model (Dataset 2) vs United States Avg

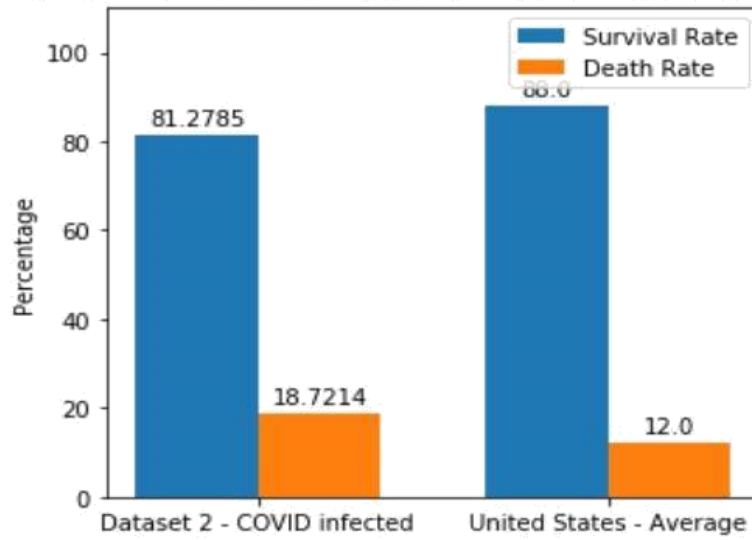


Figure 2 - Comparison of Death rate prediction of our model vs the United States closed cases

Comparison of Death rates from our model (Dataset 2 - Normal) vs United States Avg

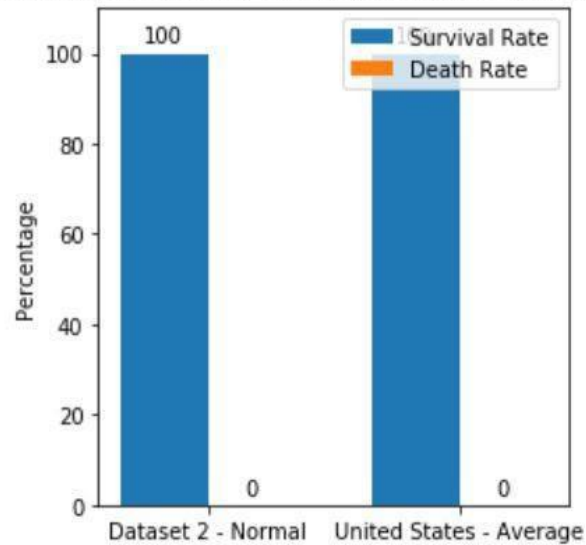


Figure 3 - Comparison of Death rate prediction of our model vs the United States case for *Non-COVID-19* Infected patients

Comparison of Death rates from our model Pneumonia Infected vs US Average

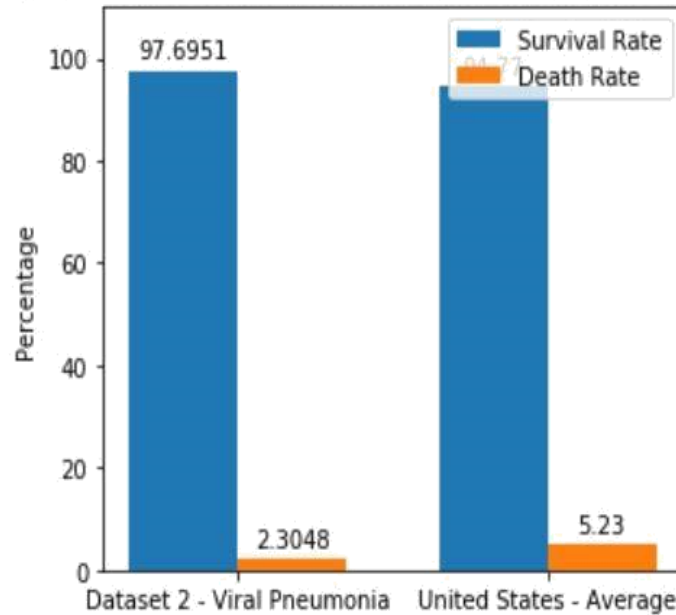


Figure 4 - Comparison of Death rate prediction of our model vs the United States closed cases for Viral Pneumonia Infected patients

CNN's model summary and architecture

Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d_11 (MaxPooling)	(None, 31, 31, 32)	0
conv2d_12 (Conv2D)	(None, 29, 29, 32)	9248
max_pooling2d_12 (MaxPooling)	(None, 14, 14, 32)	0
flatten_6 (Flatten)	(None, 6272)	0
dense_11 (Dense)	(None, 128)	802944
dense_12 (Dense)	(None, 2)	258
Total params: 813,346		
Trainable params: 813,346		
Non-trainable params: 0		

Figure 5- Architecture of the CNN model

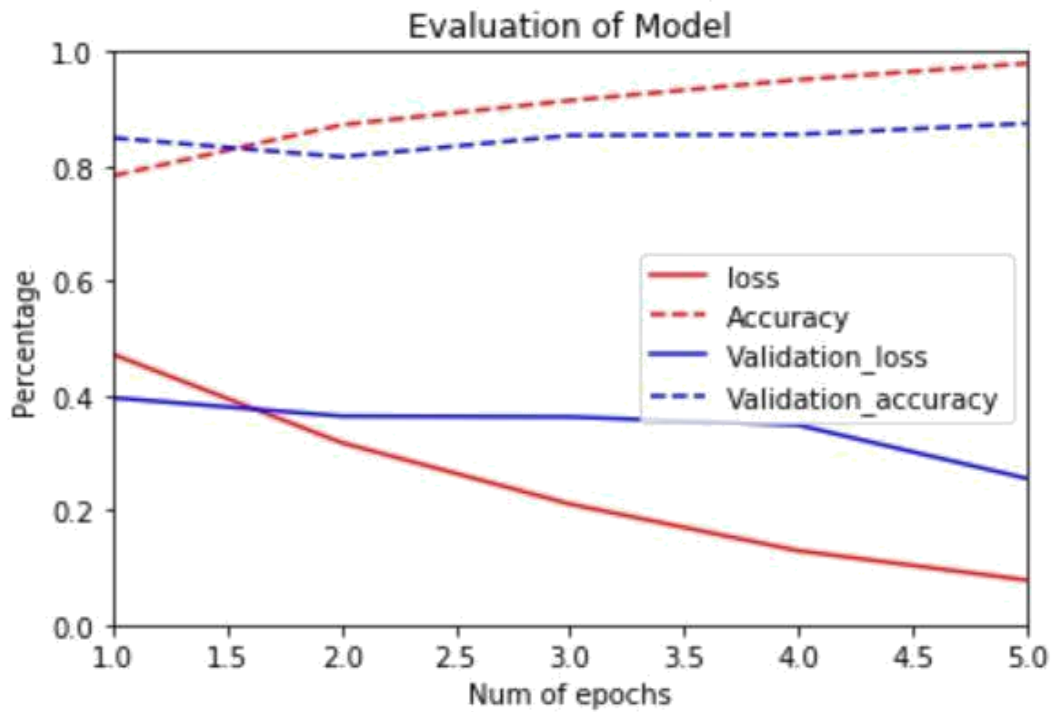


Figure 6- Evaluation of the CNN model

GAN ARCHITECTURE AND SUMMARY

Generative Model Summary

Model: "sequential_2"

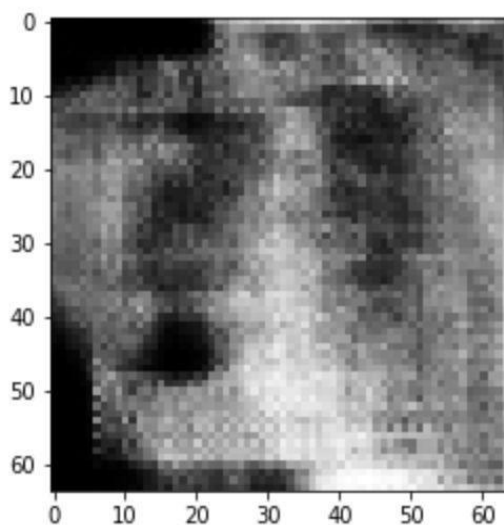
Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 32768)	3309568
leaky_re_lu_3 (LeakyReLU)	(None, 32768)	0
reshape_1 (Reshape)	(None, 16, 16, 128)	0
conv2d_transpose_1 (Conv2DTr	(None, 32, 32, 128)	262272
leaky_re_lu_4 (LeakyReLU)	(None, 32, 32, 128)	0
conv2d_transpose_2 (Conv2DTr	(None, 64, 64, 128)	262272
leaky_re_lu_5 (LeakyReLU)	(None, 64, 64, 128)	0
conv2d_3 (Conv2D)	(None, 64, 64, 1)	32769
Total params: 3,866,881		
Trainable params: 3,866,881		
Non-trainable params: 0		

Discriminative Model Summary

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 64)	640
leaky_re_lu_1 (LeakyReLU)	(None, 32, 32, 64)	0
dropout_1 (Dropout)	(None, 32, 32, 64)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	36928
leaky_re_lu_2 (LeakyReLU)	(None, 16, 16, 64)	0
dropout_2 (Dropout)	(None, 16, 16, 64)	0
flatten_1 (Flatten)	(None, 16384)	0
dense_1 (Dense)	(None, 1)	16385
Total params: 107,906		
Trainable params: 53,953		
Non-trainable params: 53,953		

The following X-Ray image is artificially generated using a GAN. The GAN network was trained on a local CPU with 50 iterations. We can drastically improve the quality of generated images given computation power and time.



Conclusion

Our model shows enough promise to be considered significant. The real-world death rate is around 12 percent, and our model predicts 19. The model also predicts ***zero mortality*** for ***normal X-ray images***. The skewness can be accounted for numerous factors, such as a biased sample size, or lack of training data. Such a close correlation between chest X-ray and mortality is an interesting endeavor to pursue and our investigation performs the task of showing noteworthy indication.

Next Steps:

The next step in our project would be to use trained and tuned GAN to further generate the training data. The generation of labeled training data would significantly increase the accuracy of our model, and we will be able to test it widely across different datasets.

